

Assignment 3: Function Approximation with Neural Network and Backpropagation

1905095 - Md Raihan Sobhan

November 8, 2024

1 Introduction

In this assignment, we had to implement a Feed-Forward Neural Network (FNN) from scratch and apply our FNN to classify apparel.

2 Instructions for Running the Code

To reproduce the results in this report, follow these steps:

1. Set up a Python virtual environment and activate it:

```
python3 -m venv myenv  
source myenv/bin/activate
```

2. Install the required packages:

```
pip install --upgrade pip  
pip install numpy matplotlib scikit-learn pillow  
pip install torch torchvision torchaudio --index-url https://download.pytorch.org/w
```

3. Select Kernel - myenv (Python 3.12.5) on VS Code.
4. Run the main Jupyter Notebook File 1905095.ipynb by executing: Run All on VS Code.

This will train the neural network on the FashionMNIST dataset, perform evaluation, and save the best model in pickle.

3 Training Results: Learning Rate and Model Comparisons

We experimented with 12 different configurations, combining four learning rates (0.005, 0.001, 0.0005, 0.0001) and three hidden layer sizes (64, 128, 256). For each combination, we recorded training loss, validation loss, accuracy, and Macro-F1 score over epochs.

3.1 Learning Rate and Model Performance

Each combination of learning rate and hidden layer size is summarized below. We present graphs showing loss, accuracy, and Macro-F1 score over epochs for each combination, followed by the confusion matrix.

3.1.1 Learning Rate: 0.005, Hidden Layer Size: 64

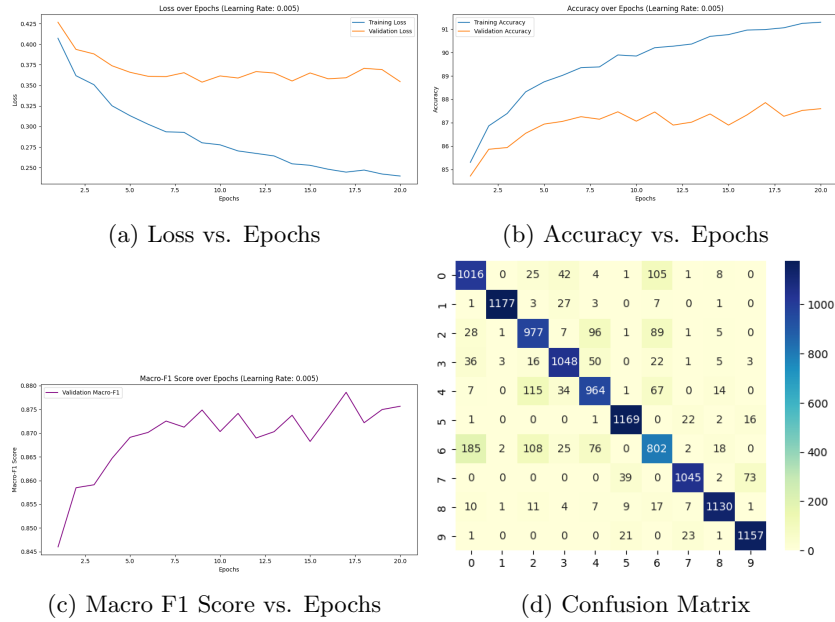


Figure 1: Metrics for Learning Rate 0.005, Hidden Layer Size 64

3.1.2 Learning Rate: 0.005, Hidden Layer Size: 128

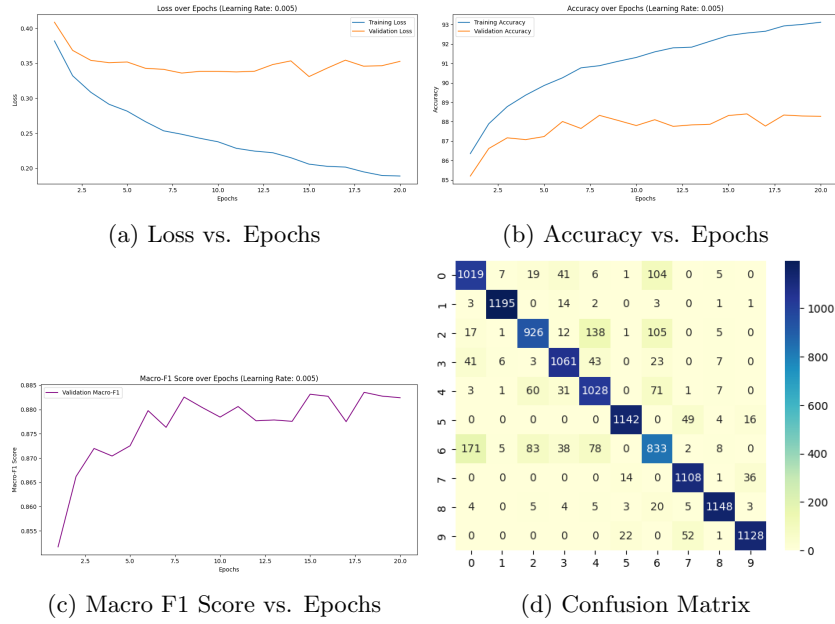


Figure 2: Metrics for Learning Rate 0.005, Hidden Layer Size 128

3.1.3 Learning Rate: 0.005, Hidden Layer Size: 256

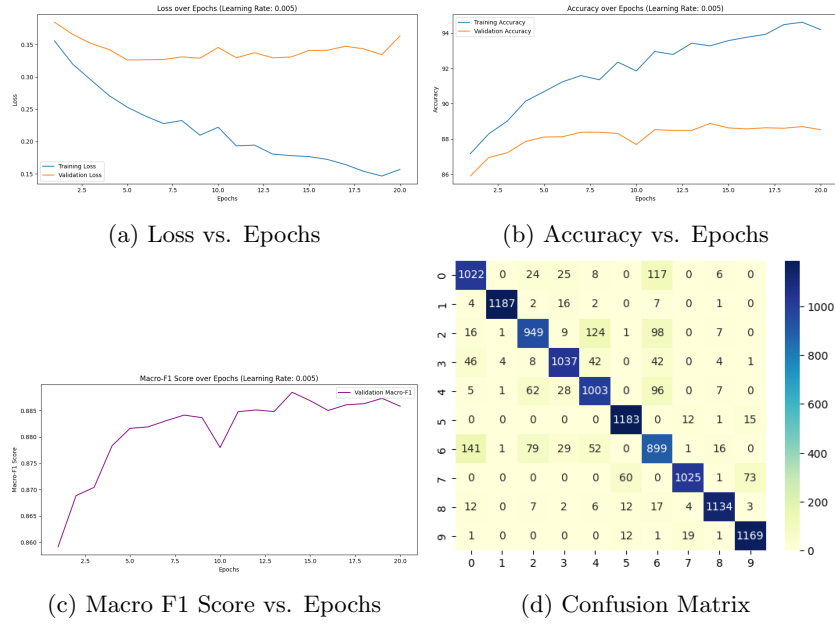


Figure 3: Metrics for Learning Rate 0.005, Hidden Layer Size 256

3.1.4 Learning Rate: 0.001, Hidden Layer Size: 64

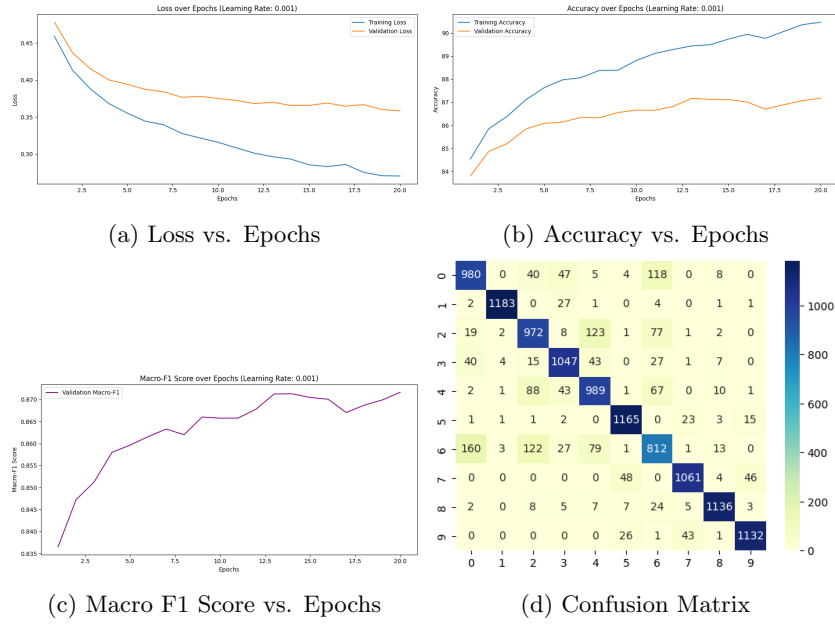


Figure 4: Metrics for Learning Rate 0.001, Hidden Layer Size 64

3.1.5 Learning Rate: 0.001, Hidden Layer Size: 128

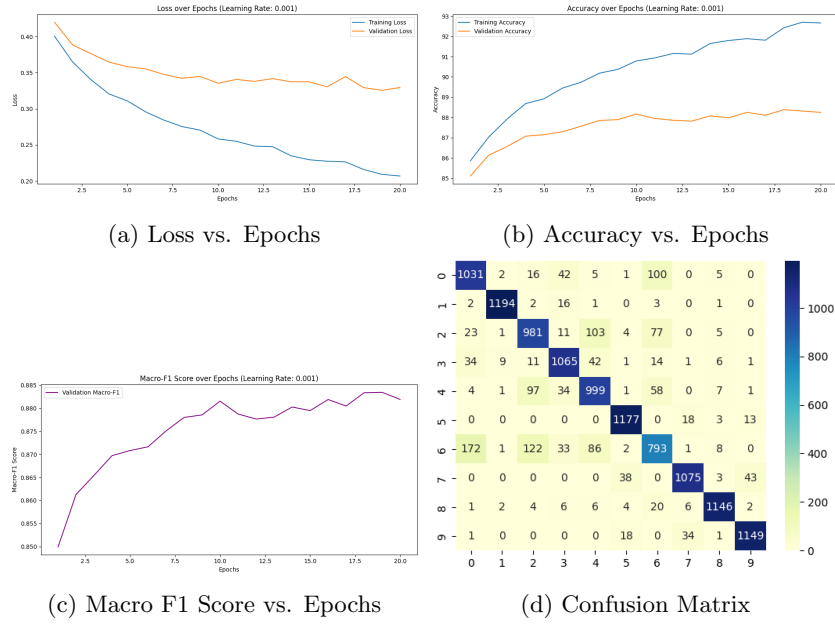


Figure 5: Metrics for Learning Rate 0.001, Hidden Layer Size 128

3.1.6 Learning Rate: 0.001, Hidden Layer Size: 256

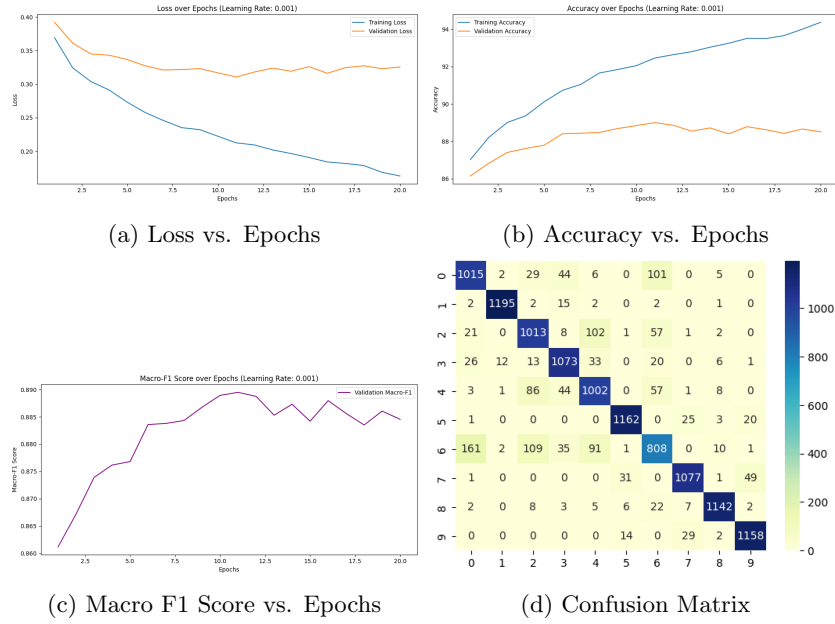


Figure 6: Metrics for Learning Rate 0.001, Hidden Layer Size 256

3.1.7 Learning Rate: 0.0005, Hidden Layer Size: 64

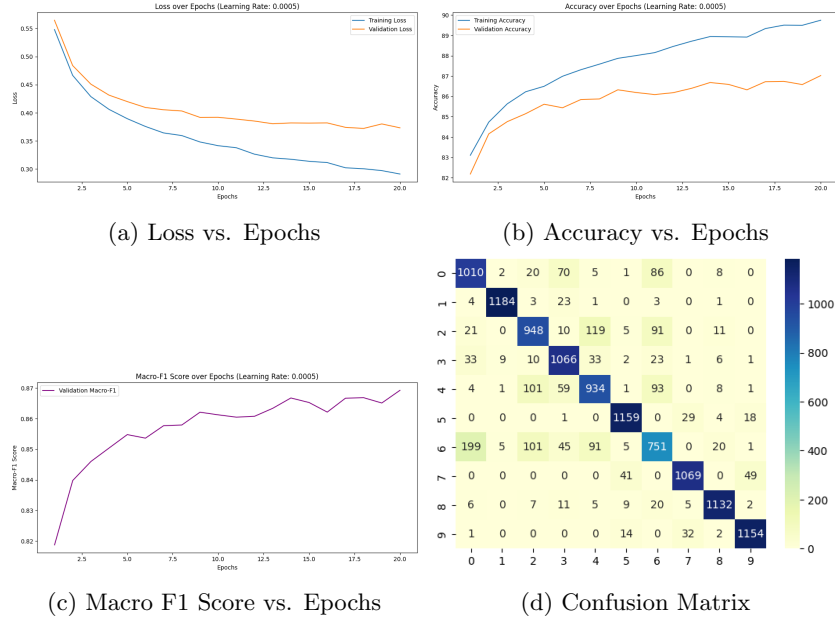


Figure 7: Metrics for Learning Rate 0.0005, Hidden Layer Size 64

3.1.8 Learning Rate: 0.0005, Hidden Layer Size: 128

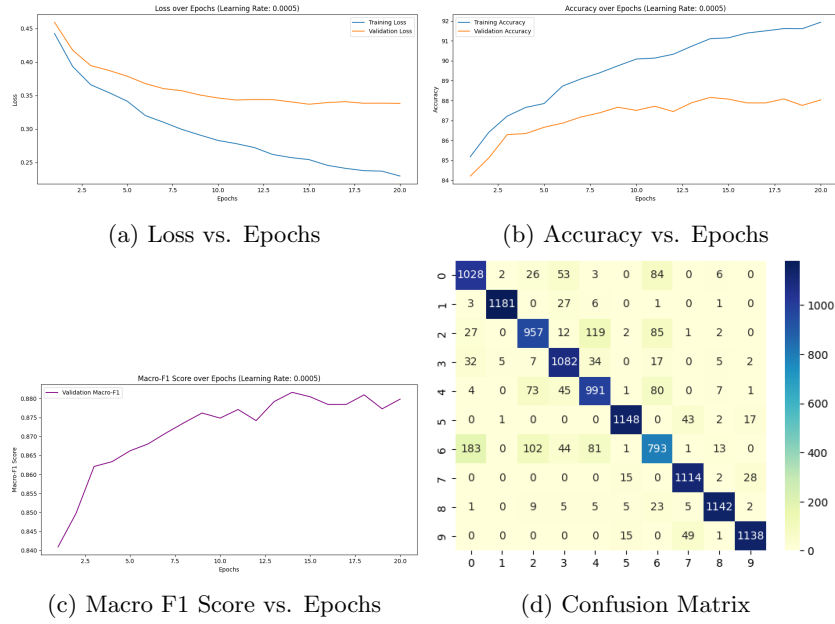


Figure 8: Metrics for Learning Rate 0.0005, Hidden Layer Size 128

3.1.9 Learning Rate: 0.0005, Hidden Layer Size: 256

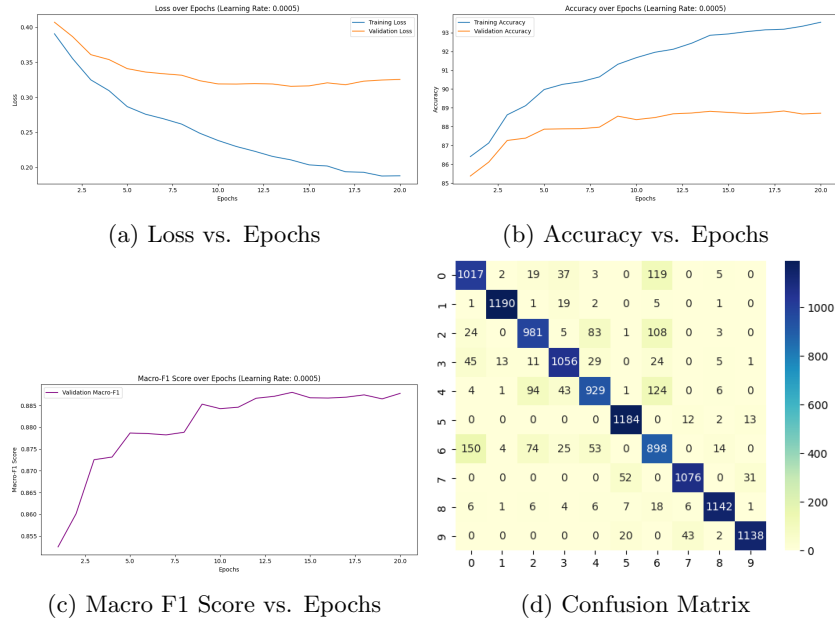


Figure 9: Metrics for Learning Rate 0.0005, Hidden Layer Size 256

3.1.10 Learning Rate: 0.0001, Hidden Layer Size: 64

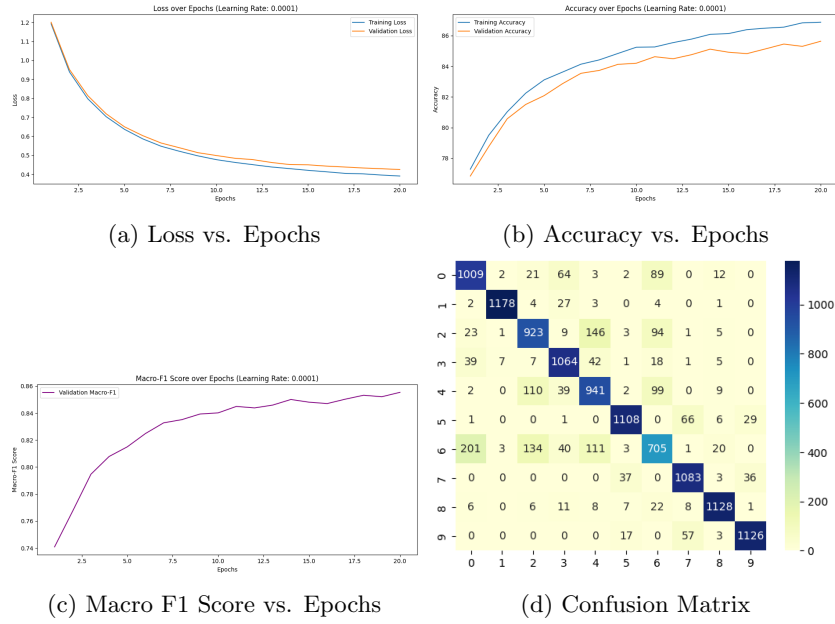


Figure 10: Metrics for Learning Rate 0.0001, Hidden Layer Size 64

3.1.11 Learning Rate: 0.0001, Hidden Layer Size: 128

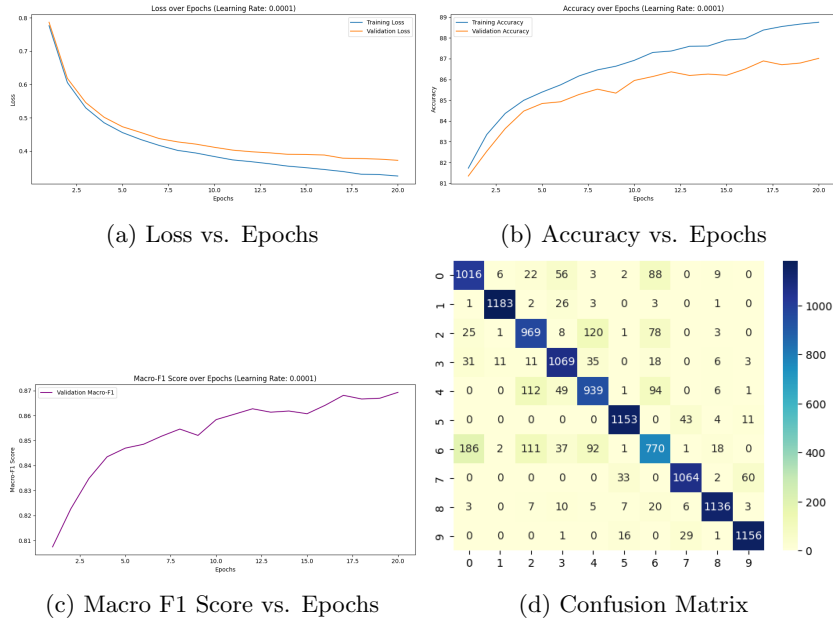


Figure 11: Metrics for Learning Rate 0.0001, Hidden Layer Size 128

3.1.12 Learning Rate: 0.0001, Hidden Layer Size: 256

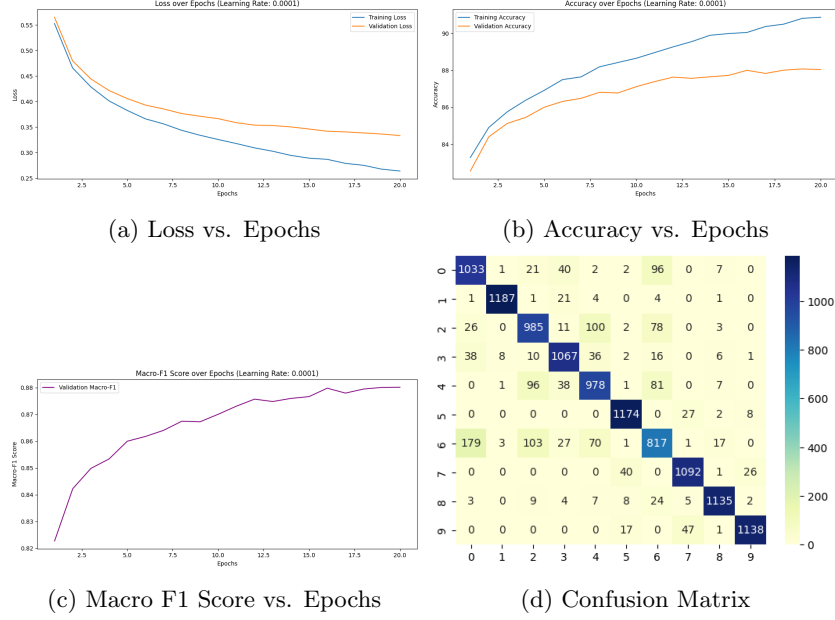


Figure 12: Metrics for Learning Rate 0.0001, Hidden Layer Size 256

4 Comparison and Selection of the Best Model

From the experiments, we identified the best model based on the highest validation Macro-F1 score. Table 1 summarizes the final metrics for each configuration.

Table 1: Validation Metrics for Each Model Configuration

LR	Hidden Layer	Train Loss	Val Loss	Train Acc	Val Acc	Val Macro-F1
0.005	64	0.2343	0.3603	91.35%	87.17%	0.8714
0.005	128	0.1874	0.3494	93.11%	87.94%	0.8791
0.005	256	0.1558	0.3437	94.22%	88.52%	0.8857
0.001	64	0.2693	0.3616	90.41%	87.04%	0.8701
0.001	128	0.2085	0.3382	92.61%	88.36%	0.8837
0.001	256	0.1606	0.3254	94.39%	88.93%	0.8890
0.0005	64	0.2938	0.3706	89.76%	87.07%	0.8705
0.0005	128	0.2292	0.3380	92.09%	88.15%	0.8812
0.0005	256	0.1864	0.3245	93.35%	88.38%	0.8835
0.0001	64	0.3901	0.4288	86.90%	85.32%	0.8526
0.0001	128	0.3194	0.3709	88.99%	87.20%	0.8714
0.0001	256	0.2649	0.3361	90.92%	88.08%	0.8798

5 Test Performance of the Best Model

The model with the highest validation Macro-F1 score was chosen as the best model (LR = 0.001 and Hidden Layer Size 256). We evaluated it on an independent test set and recorded the following metrics:

- **Test Accuracy:** 89.07
- **Test Macro-F1 Score:** 0.8900029201513565




Figure 13: Confusion Matrix for Best Model on Test Set

6 Conclusion

This report presents the results of training a neural network on the FashionM-NIST dataset, evaluating performance with 12 configurations of learning rate and hidden layer size. The optimal model achieved a validation Macro-F1 score of 0.89 and also achieved a strong test accuracy of 89.07